

MATH3871 - Assignment

Yue Yu - z5359109

T3 2024

Part I: Frequentist analysis

1).

```
1 wine = read.csv("winequality-red.csv")
2 anyNA(wine)
3 wine = na.omit(wine)
```

2).

```
1 wine$good = ifelse(wine$quality >= 6.5, 1, 0)
```

3).

```
1 model = glm(good ~ . - quality, data = wine, family = binomial)
2 mleest = coef(model)
```

Part II: Bayesian analysis

1).

We start from the density function of logistic regression model:

$$P(y_i = 1|\beta, \mathbf{x}_i) = \frac{\exp(\mathbf{x}_i\beta)}{1 + \exp(\mathbf{x}_i\beta)}.$$

$$\begin{aligned} P(y_i|\beta, \mathbf{x}_i) &= P(y_i = 1|\beta, \mathbf{x}_i)^{y_i} [1 - P(y_i = 1|\beta, \mathbf{x}_i)]^{(1-y_i)} \\ &= \left[\frac{\exp(\mathbf{x}_i\beta)}{1 + \exp(\mathbf{x}_i\beta)} \right]^{y_i} \left[\frac{1}{1 + \exp(\mathbf{x}_i\beta)} \right]^{(1-y_i)} \\ &= \frac{\exp(y_i\mathbf{x}_i\beta)}{1 + \exp(\mathbf{x}_i\beta)}. \end{aligned}$$

Then we compute the likelihood function $L(y|\beta, \mathbf{x})$

$$L(y|\beta, \mathbf{x}) = \prod_{i=1}^n P(y_i|\beta, \mathbf{x}_i).$$

Hence the log-likelihood being

$$\begin{aligned} \log L(y|\beta, \mathbf{x}) &= \log \prod_{i=1}^n P(y_i|\beta, \mathbf{x}_i) \\ &= \sum_{i=1}^n \log P(y_i|\beta, \mathbf{x}_i) \\ &= \sum_{i=1}^n \log \frac{\exp(y_i\mathbf{x}_i\beta)}{1 + \exp(\mathbf{x}_i\beta)} \\ &= \sum_{i=1}^n y_i\mathbf{x}_i\beta - \log [1 + \exp(\mathbf{x}_i\beta)]. \end{aligned}$$

2).

We know that the posterior is

$$\pi(\beta|y, \mathbf{x}) \propto L(y|\beta, \mathbf{x})\pi(\beta)$$

Hence the log-posterior can be obtained by

$$\begin{aligned} \log \pi(\beta|y, \mathbf{x}) &\propto \log [L(y|\beta, \mathbf{x})\pi(\beta)] \\ &= \log L(y|\beta, \mathbf{x}) + \log \pi(\beta), \end{aligned}$$

which is a sum of log-prior and log-likelihood. Since we've derived the log-likelihood in the previous part, we now compute the log-prior. Since $\alpha = 0$, the prior is

$$\pi(\beta) = \frac{1}{(2\pi)^{(k+1)/2} |\Omega|^{1/2}} \exp \left(-\frac{1}{2} \beta^T \Omega^{-1} \beta \right).$$

Hence substitute $\Omega = 100I_{k+1}$ the log-prior is

$$\begin{aligned}\log \pi(\beta) &= -\frac{k+1}{2} \log(2\pi) - \frac{1}{2} \log |\Omega| - \frac{1}{2} \beta^T \Omega^{-1} \beta \\ &= -\frac{k+1}{2} \log(2\pi) - \frac{1}{2} \log |100I_{k+1}| - \frac{1}{2} \beta^T (100I_{k+1})^{-1} \beta \\ &= -\frac{k+1}{2} \log(2\pi) - \frac{1}{2} \log 100 - \frac{1}{200} \beta^T \beta,\end{aligned}$$

and the log-posterior is

$$\begin{aligned}\log \pi(\beta|y, \mathbf{x}) &\propto \log L(y|\beta, \mathbf{x}) + \log \pi(\beta) \\ &= \sum_{i=1}^n \{y_i \mathbf{x}_i \beta - \log [1 + \exp(\mathbf{x}_i \beta)]\} - \frac{k+1}{2} \log(2\pi) - \frac{1}{2} \log 100 - \frac{1}{200} \beta^T \beta.\end{aligned}$$

```

1 lpost.LR <- function(beta,X,y) {
2
3   # log-likelihood
4   x_beta = X %*% beta
5   log_L = sum(y * x_beta - log(1 + exp(x_beta)))
6
7   # log-prior
8   k = length(beta) - 1
9   log_prior = - (k + 1) / 2 * log(2 * pi) - 0.5 * (k + 1) * log(100) - 1/200 * t(beta
10   ) %*% beta
11
12   # log posterior
13   log_post = log_L + log_prior
14   return(log_post)
15 }
```

3).

Note that for numeric stability, we usually compute

$$\log \left(\frac{\pi(\theta^*)q(\theta^*, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta^*)} \right) = \log \pi(\theta^*) + \log q(\theta^*, \theta_t) - \log \pi(\theta_t) - \log q(\theta_t, \theta^*),$$

and compute $\log u$ with $u \sim U(0, 1)$. Since in this case our proposal distribution

$$q(\beta|\beta_t) = N_{k+1}(\beta_t, \Sigma)$$

is symmetrical, we can conclude that

$$q(\theta^*, \theta_t) = q(\theta_t, \theta^*).$$

Hence,

$$\log \left(\frac{\pi(\theta^*)q(\theta^*, \theta_t)}{\pi(\theta_t)q(\theta_t, \theta^*)} \right) = \log \pi(\theta^*) - \log \pi(\theta_t),$$

which is used in the code below as the acceptance criteria.

```

1 mhmcmc <- function(y, X, B, nsims, Sigma) {
2
3   k = length(B)
4   beta_mat = matrix(NA, nrow = nsims, ncol = k)
5   beta_mat[1, ] = B
6   acc = 0
7
8   for (t in 2:nsims) {
9
10    # Curr and proposed Beta
11    beta_curr = beta_mat[t - 1, ]
12    beta_star = mvrnorm(1, mu = beta_curr, Sigma = Sigma)
13
14    log_post_curr = lpost.LR(beta_curr, X, y)
15    log_post_star = lpost.LR(beta_star, X, y)
16    accprob = exp(log_post_star - log_post_curr)
17
18    if (log(runif(1)) < log_post_star - log_post_curr) {
19      beta_mat[t, ] = beta_star
20      acc = acc + 1
21    } else {
22      beta_mat[t, ] = beta_curr
23    }
24  }
25  mhout = list("beta_mat"=beta_mat, "accprob"=accprob, "beta_star" = beta_star, "acc"
26              =acc)
27  return(mhout)
28 }

```

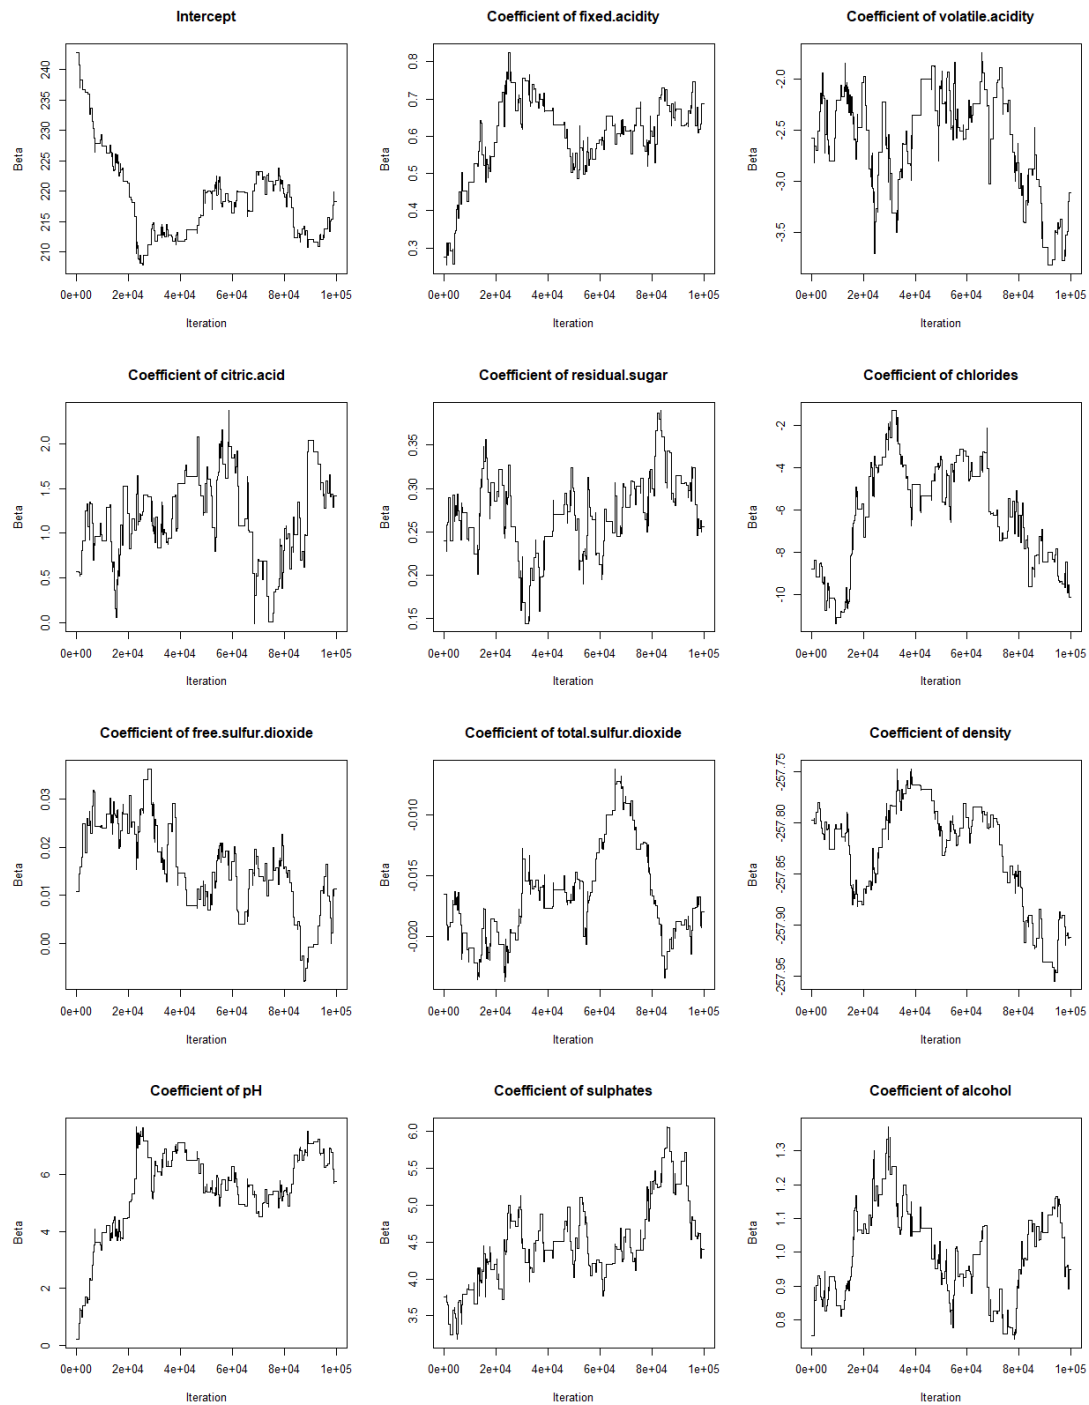
4).

```

1 # Covariance for proposal:
2 Sigma = diag(c(1100, 0.0015, 0.04, 0.05, 0.0005, 0.4, 0.00001, 0.000001, 0.0001,
3               0.09, 0.03, 0.002)) #Do not modify
4 nsims = 1e5
5 X = cbind(Intercept = 1, as.matrix(wine[, !names(wine) %in% c("good", "quality")]))
6 y = wine$good
7 mhout1 = mhmcmc(y, X, mleest, nsims, Sigma)

```

a).



```
1 par(mfrow=c(4,3))
2 for (j in 1:length(mleest)) {
3   name = if (j==1) "Intercept" else paste("Coefficient of", colnames(wine)[j-1])
4
5   plot(mhout1$beta_mat[, j], type = "l", main = name, xlab = "Iteration", ylab = "
      Beta")
6 }
```

b).

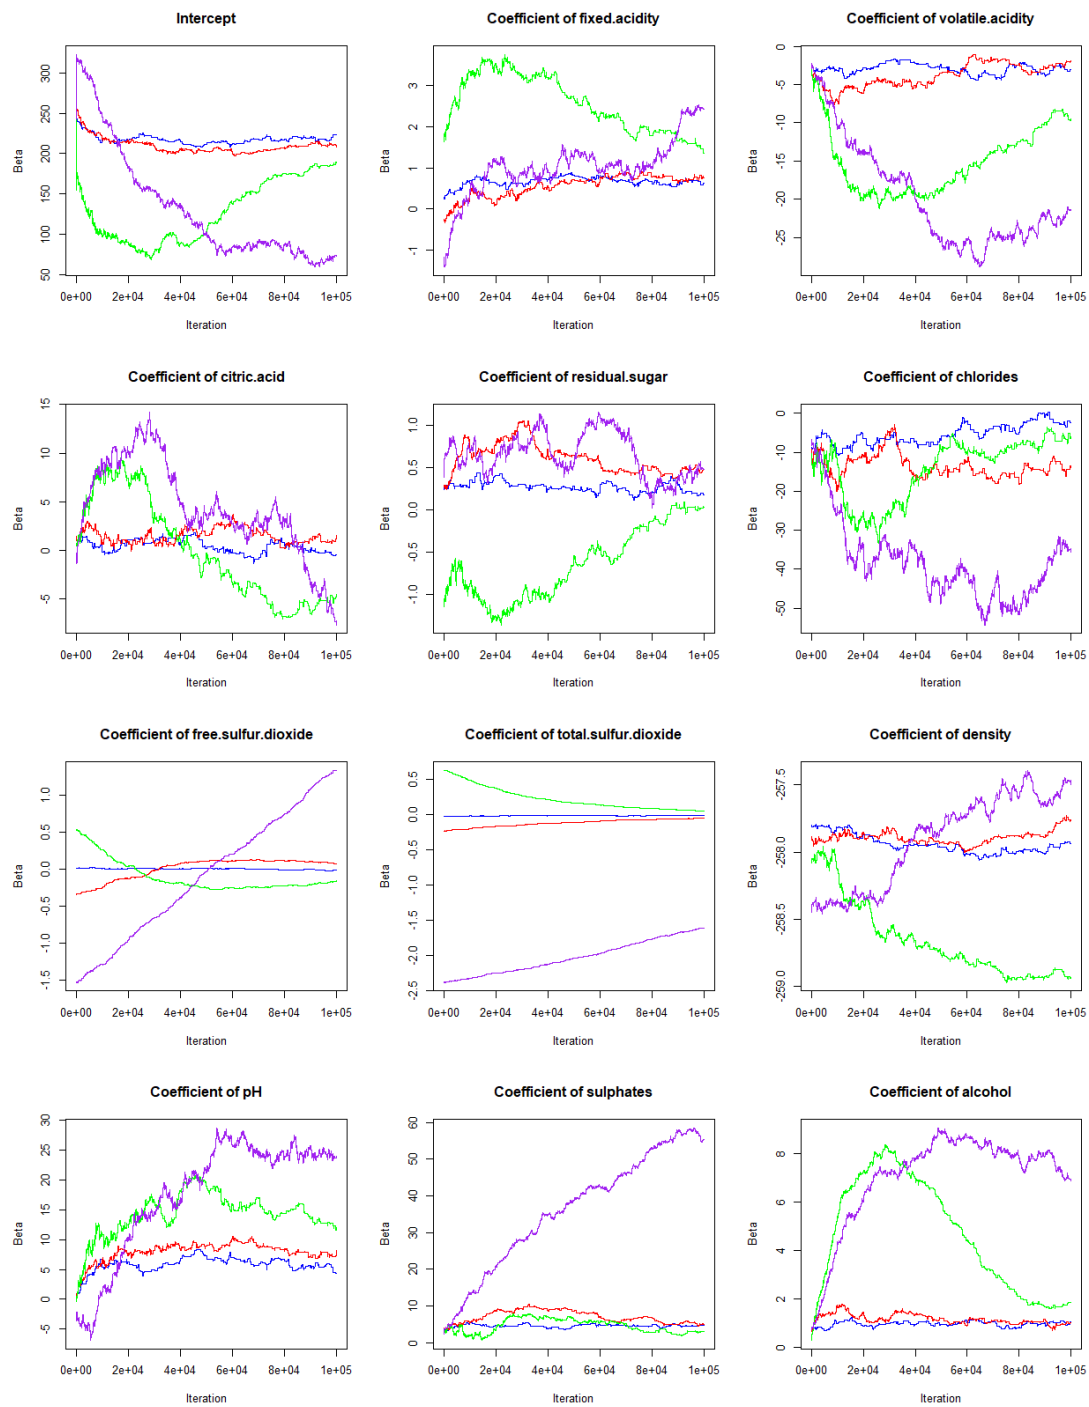
```
1 mh4acc = mhout1$acc / nsims
```

$$mh4acc = 0.00232 = 0.2\%.$$

c).

From the graph, we can see there exist no clear sign of convergence for the coefficients, hence the proposed distribution is not very suitable. Also the acceptance is very low at 0.2%, and the initialization is far away from the MLE, hence the initialization is not suitable either. This conclusion aligns with the fact that the rate of acceptance is low.

5).



The lines above have the initialization:

```

blue : mleest,
red : mleest + rnorm(length(mleest), 0, 0.5),
green : mleest + rnorm(length(mleest), 0, 1),
purple : mleest + rnorm(length(mleest), 0, 1.5).

```

```

1 B_init2 = mleest
2 B_init3 = mleest + rnorm(length(mleest), 0, 0.5)
3 B_init4 = mleest + rnorm(length(mleest), 0, 1)
4 B_init5 = mleest + rnorm(length(mleest), 0, 1.5)
5
6 mhout2 = mcmc(y, X, B_init2, nsims, Sigma)
7 mhout3 = mcmc(y, X, B_init3, nsims, Sigma)
8 mhout4 = mcmc(y, X, B_init4, nsims, Sigma)
9 mhout5 = mcmc(y, X, B_init5, nsims, Sigma)
10
11 par(mfrow = c(4, 3))
12
13 for (j in 1:length(mleest)) {
14   name = if (j==1) "Intercept" else paste("Coefficient of", colnames(wine)[j-1])
15   plot(mhout2$beta_mat[, j], type = "l", col = "blue", main = name, ylim = range(
16     mhout2$beta_mat[, j], mhout3$beta_mat[, j], mhout4$beta_mat[, j], mhout5$beta_mat
17     [, j]), xlab = "Iteration", ylab = "Beta")
18   lines(mhout3$beta_mat[, j], col = "red")
19   lines(mhout4$beta_mat[, j], col = "green")
20   lines(mhout5$beta_mat[, j], col = "purple")
21 }

```

6).

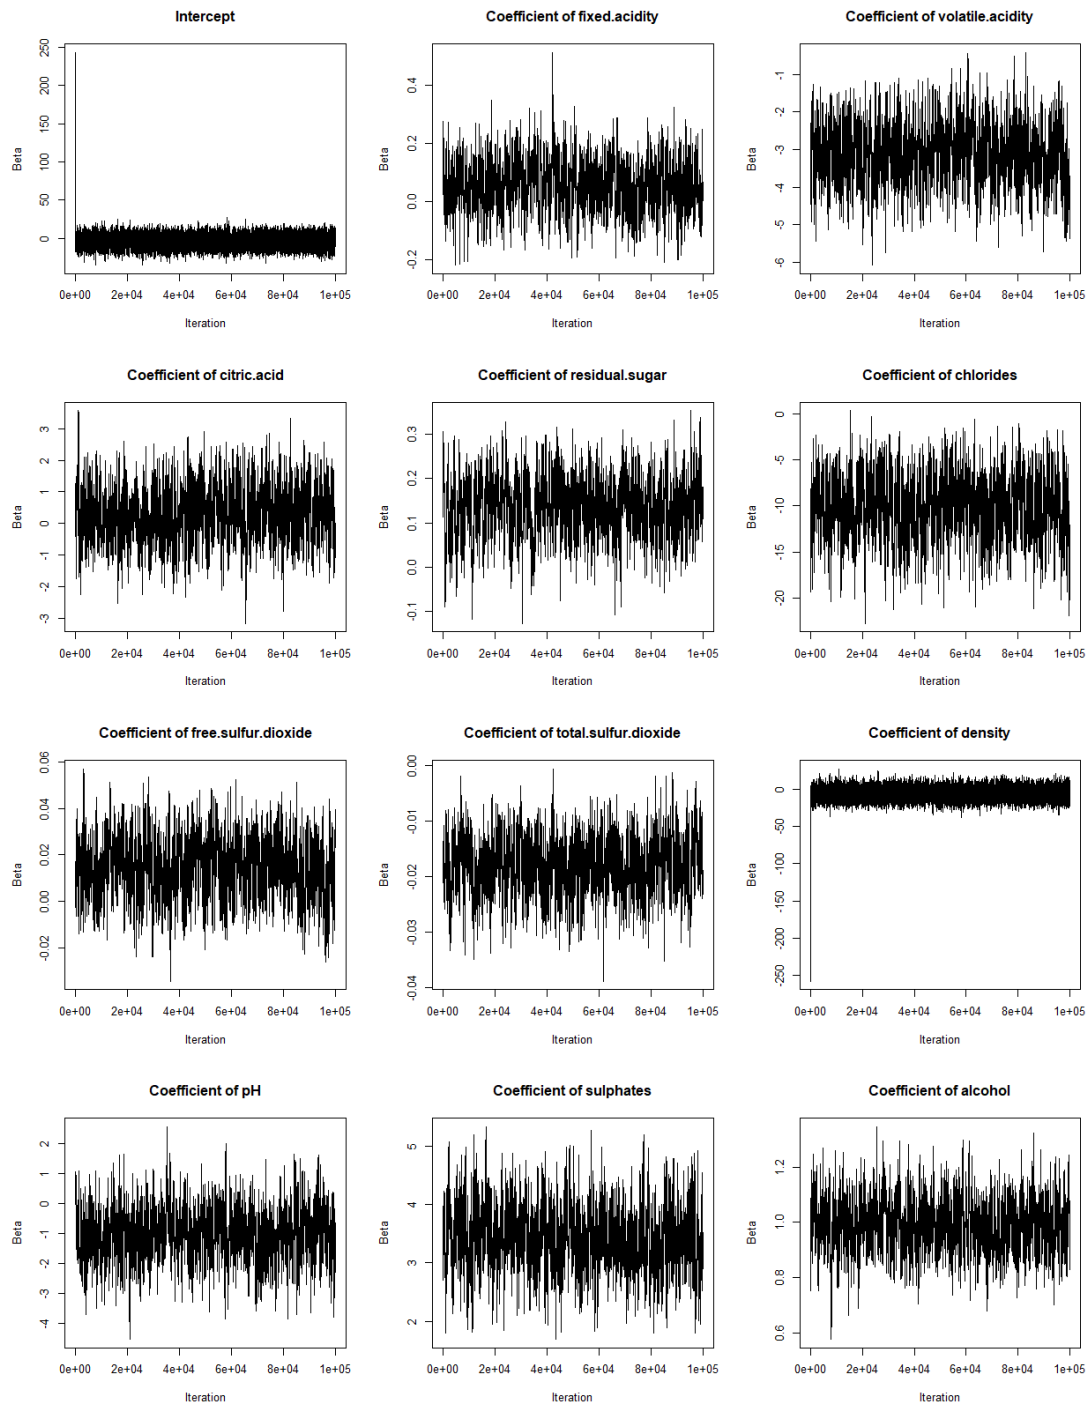
After tuning the the model, the acceptance rate is at 18.5%

```

1 cov_mle = vcov(model)
2 Sigma = cov_mle * 0.1
3 mhout6 = mcmc(y, X, mleest, nsims, Sigma)
4
5 #calculate the acceptance rate
6 mh6acc = mhout6$acc / nsims

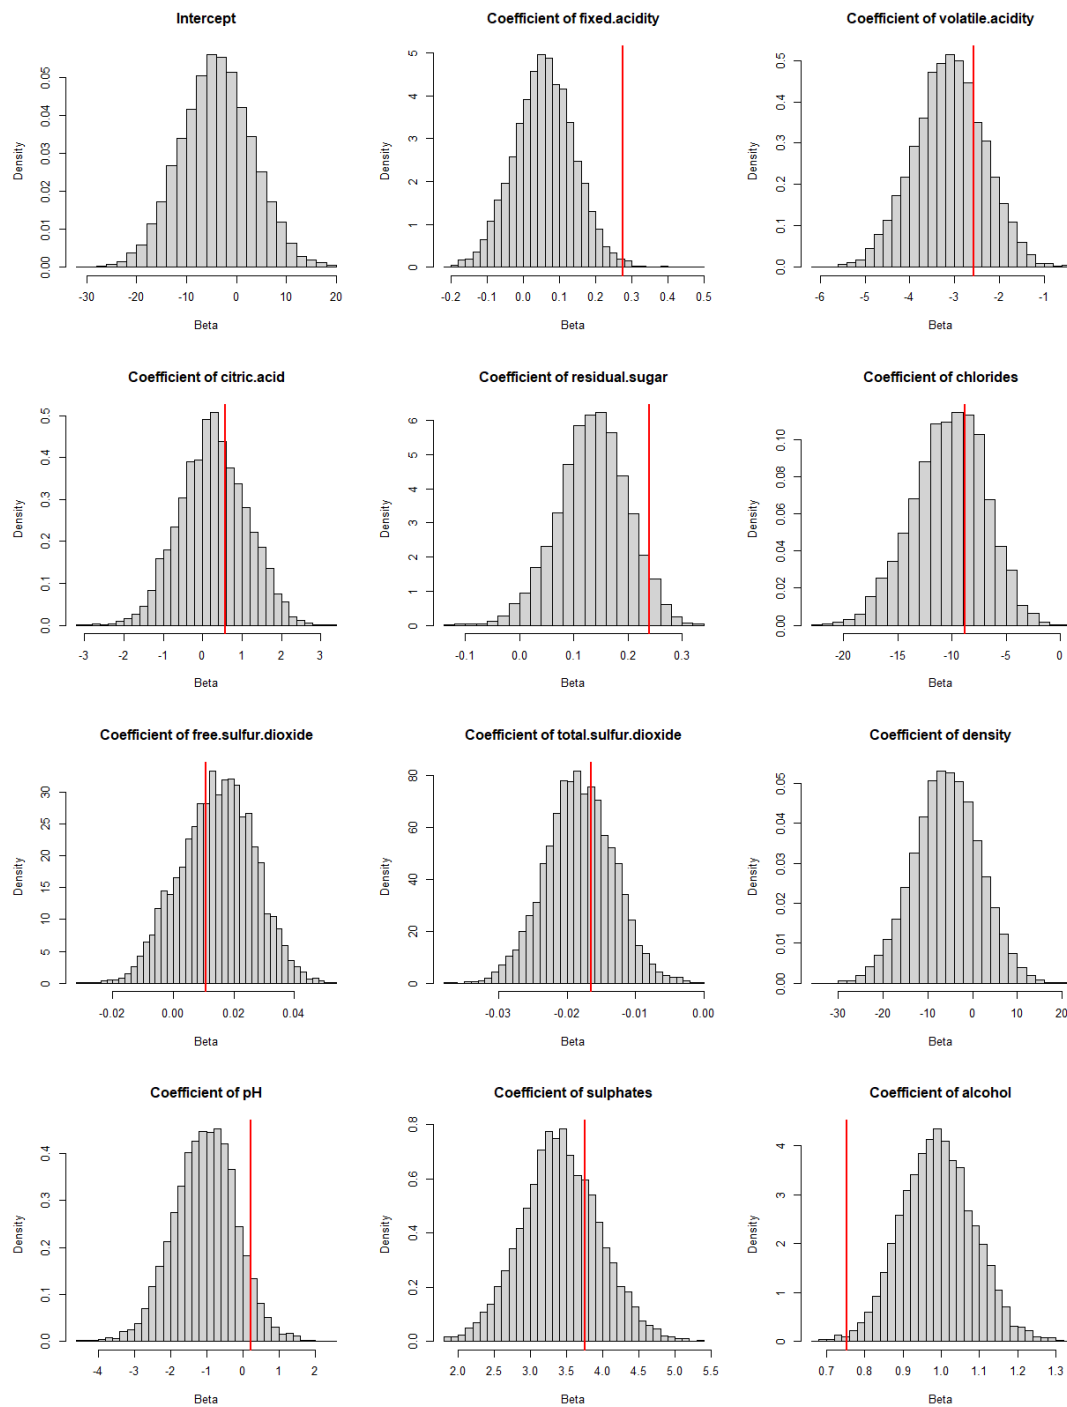
```

a).



```
1 par(mfrow=c(4,3))
2 for (j in 1:length(mleest)) {
3   name = if (j==1) "Intercept" else paste("Coefficient of", colnames(wine)[j-1])
4   plot(mhout6$beta_mat[, j], type = "l", main = name, xlab = "Iteration", ylab = "
   Beta")
5 }
```

b).



```
1 # First 10% as burn-in and keep every 10th to avoid autocorrection
2 burn_in = 0.1 * nsims
3 thin = 10
4
5 remaining = mhout6$beta_mat[(burn_in + 1):nsims, ]
6 thinned = remaining[seq(1, nrow(remaining), by = thin), ]
7 mleest[1]
8 par(mfrow = c(4, 3))
9 for (j in 1:ncol(thinned)) {
10   name = if (j==1) "Intercept" else paste("Coefficient of", colnames(wine)[j-1])
11   hist(thinned[, j], probability = TRUE, main = name, xlab = "Beta", breaks = 30)
12   abline(v = mleest[j], col = "red", lwd = 2)
13 }
```

c).

From the graph, we can see the tuned graph shows clear sign of convergence, and the rate of acceptance is raised to around 18.5%. Also to solve the burn-in issue, we are discarding the first 10% of the sample. Also to avoid auto-correction, we employed **thinning** by collecting the every 10th sample. Hence this proposal is much better suited for this problem.